

GWAtoolbox
An R package for the fast processing of data from
Genome-Wide Association Studies

Christian Fuchsberger Daniel Taliun Cristian Pattaro

October 28, 2011

Contents

1	Introduction	3
2	Installation	3
2.1	Windows	4
2.2	Unix	4
2.3	Mac OS X	5
3	The Quality Control Workflow	6
4	GWAS Data Files	8
5	The Input Script	8
5.1	Listing the Input Data Files	9
5.2	Describing the Input Data Columns	9
5.2.1	Field Separator	9
5.2.2	Missing Values	10
5.2.3	Column Names	10
5.2.4	Case Sensitivity	12
5.3	Specifying Data Filters	12
5.3.1	Implausible Values Filters	12
5.3.2	High Quality Filters	13
5.3.3	Plotting Filters	14
5.4	Controlling The Output	15
5.4.1	Output File Names	15
5.4.2	Verbosity Level For Graphical Output	15
5.4.3	Number And Content Of Plots	15
6	The Output Files	16
7	Example	18
8	Between-study comparisons	20
8.1	Comparing skewness and kurtosis of effect size distribution	20
8.2	Precision of the effect estimates by sample size	22
9	Additional Tools	23
	Index	25

1 Introduction

GWAtoolbox is an *R* package for processing data originated from Genome-Wide Association Studies (GWAS). GWAS have become increasingly popular in the last years, leading to the discovery of hundreds of common genetic variants affecting the risk of diseases (such as diabetes, hypertension, chronic kidney disease, etc.) or the level of quantitative biological parameters.

Results from GWAS typically consist of large files where, for each single nucleotide polymorphism (SNP), statistics related to the association between the SNP and the studied trait are stored. The number of SNPs which is currently being analyzed in most GWAS is in excess of 2.5 Million and is expected to increase rapidly. For each individual SNP, the minimal information stored consists of the SNP unique name, chromosomal position, genotype (reference and non-reference alleles), frequency of the reference allele, SNP effect size, its standard error, and p-value. Additional information such as minor allele frequency (MAF) and imputation quality are often provided. As a consequence, the typical dimension of GWAS result files is of >2.5 Million rows by >10 columns, for a total file size which is often larger than 300 Mbytes.

With the aim of detecting common or less common genetic variants with modest effects, it is now common practice to pool results from individual studies into meta-analysis efforts which not rarely involve dozens of studies. In these consortia initiatives, each individual study contributes from one to several files, either because multiple traits are being analyzed or because different analyses on the same trait are needed. Consequently, data analysts working in these consortia have to deal with a massive amount of files which need to be quality controlled to avoid problems during the meta-analysis process. As a result of the quality control (QC) process, some files could be found to be corrupted or erroneous so that new data upload is needed from individual studies. In this way, the loop between the consortium and the individual study analyst originates multiple file checks, until a satisfactory data quality is achieved.

When working with such large datasets in *R*, simple operations such as uploading the GWAS files into the *R* working space, file management, and data plotting, can take a considerable amount of time, and a systematic QC of hundreds of GWAS files can be unfeasible or may require several weeks.

The *GWAtoolbox* provides a set of instruments to simplify the data handling in the framework of meta-analyses of GWA data. The function *gwasqc()* is capable to process a high number of GWAS data files in a single run, and producing several QC reports and figures. Routines for the between-study comparison are also provided to check systematic difference between files. In addition, the package contains annotation and graphical tools to assist the result interpretation.

2 Installation

The *GWAtoolbox* package can be downloaded from <http://www.eurac.edu/GWAtoolbox.html>. It requires *R* version 2.13.1 or higher. The installation procedure depends on the host operating system and user privileges. In the following, detailed installation instructions for a wide range of settings are provided.

2.1 Windows

GWAtoolbox for Windows is distributed in compiled binary format. Installation:

1. Download the latest version of the package: *GWAtoolbox_X.Y.Z.zip*.
2. Start the R program.
- 3a. If you have administrator privileges, you can install the package to the main R library:

- i. Execute the command:

```
install.packages("path/to/GWAtoolbox_X.Y.Z.zip",  
repos=NULL)
```

where *path/to* is the directory where the package was downloaded.

- ii. Load the package in R with the command:

```
library(GWAtoolbox)
```

- 3b. If you do NOT have administrator privileges:

- i. Execute the command:

```
install.packages("path/to/GWAtoolbox_X.Y.Z.zip",  
lib="path/to/install/directory",  
repos=NULL)
```

where *path/to* is the directory where the package was downloaded, and *path/to/install/directory* is the path with your install directory.

- ii. Load the package in R with command:

```
library(GWAtoolbox, lib.loc =  
"path/to/install/directory")
```

2.2 Unix

GWAtoolbox for Unix is distributed in source format and, therefore, it needs to be compiled on the user machine. This requires the C/C++ compilers to be available on the system. If the requirement is fulfilled, the package can be installed as follows:

1. Download the latest package version *GWAtoolbox_X.Y.Z.tar.gz*.
- 2a. If you have administrator privileges, you can install packages to the main R library:

- i. In the Unix shell execute the command:

```
R CMD INSTALL path/to/GWAtoolbox_X.Y.Z.tar.gz
```

where *path/to* is the directory where the package was downloaded.

- ii. Start the R program and load the package with the command:

```
library(GWAtoolbox)
```

- 2b. If you do NOT have administrator privileges, follow the following steps:

- i. In the Unix shell execute the single line command:

```
R CMD INSTALL path/to/GWAtoolbox_X.Y.Z.tar.gz
-l path/to/install/directory
```

where `path/to` is the directory where the package was downloaded, and `path/to/install/directory` is the path with your install directory.

- ii. Start the R program and load the package with the command:

```
library(GWAtoolbox, lib.loc="path/to/install/directory")
```

2.3 Mac OS X

GWAtoolbox for Mac OS X is distributed in compiled binary format. The following steps describe the installation procedure:

1. Download the latest package version *GWAtoolbox_X.Y.Z.tar.gz*.

- 2a. To install from the Mac OS X shell (*Terminal*):

- i. If you have administrator privileges, you can install packages to the main R library:

- A. In the Mac OS X shell execute the command:

```
R CMD INSTALL path/to/GWAtoolbox_X.Y.Z.tar.gz
```

where `path/to` is the directory where the package was downloaded.

- B. Start the R program and load the package with the command:

```
library(GWAtoolbox)
```

- ii. If you do NOT have administrator privileges:

- A. In the Mac OS X shell execute the single line command:

```
R CMD INSTALL path/to/GWAtoolbox_X.Y.Z.tar.gz
-l path/to/install/directory
```

where `path/to` is the directory where the package was downloaded, and `path/to/install/directory` is the path with your install directory.

- B. Start the R program and load the package with the command:

```
library(GWAtoolbox, lib.loc="path/to/install/directory")
```

- 2b. To install from R:

- i. Start the R program.

- ii. If you have administrator privileges, you can install packages to the main R library:

- A. Execute the command:

```
install.packages("path/to/GWAtoolbox_X.Y.Z.tar.gz",
repos=NULL)
```

where `path/to` is the directory where the package was downloaded.

- B. Load the package in R with the command:

```
library(GWAtoolbox)
```

iii. If you do NOT have administrator privileges:

A. Execute the command:

```
install.packages("path/to/GWAtoolbox_X.Y.Z.tar.gz",  
lib="path/to/install/directory",  
repos=NULL)
```

where `path/to` is the directory where the package was downloaded, and `path/to/install/directory` is the path with your install directory.

B. Load the package in R with the command:

```
library(GWAtoolbox, lib.loc =  
"path/to/install/directory")
```

3 The Quality Control Workflow

A careful and thorough data QC should be performed before starting any meta-analysis of GWAS data, especially when many studies are involved. In this framework, we identified three complementary aspects of a good QC analysis:

1. Formal checking: control that all files that are going to be submitted to the meta-analysis fulfill some formatting guidelines, including:
 - consistency of column names with meta-analysis guidelines;
 - presence of the minimal required information;
 - data are in a format that can be analyzed (numeric, character, factor);
 - all SNP identification numbers are unique;
 - alleles are coded in letters/numbers as expected;
 - missing values are coded in a consistent way;
 - the field separator is as expected;
 - strand information is present and unequivocal;
 - the number of chromosomes and chromosome coding are as expected;
2. Quality checking: evaluating the quality of data in each single file. This includes:
 - assessing the presence of unexpected values for some of the items required for the meta-analysis (e.g.: negative p-values or negative standard errors);
 - assessing p-value inflation and analyzing p-value distribution;
 - assessing the distribution of the main summary statistics, including the effect estimates, their standard errors, genotype imputation quality, etc.

3. Global checking: identification of any systematic bias that may affect the analysis. This step is aimed to uncover studies that are systematically different from the others. This may happen when, for instance, analysts of one study forgot to log-transform the phenotype or apply the wrong model to the data.

Formal and quality checks of individual studies are performed in *GWAtoolbox* using the *gwasqc()* function. *gwasqc()* was built to address specific requirements, specifically:

1. it allows rapid file processing and reporting;
2. it eliminates routine user operations;
3. it allows multi-format reporting, including *HTML*, *CSV*, and text files.

The complete QC workflow can be summarized in four basic steps (see Figure 1):

1. collect the GWAS data files;
2. write an input script to process of all GWAS files with the *gwasqc()* function;
3. run the QC using *gwasqc()*;
4. analyze the QC results to uncover errors or inconsistencies.




Figure 1: The quality control workflow.

In the next sections we cover each of the four steps and describe the requirements for the input files and the precise content of all output files.

4 GWAS Data Files

GWAS data are usually stored as delimited text files. The first line of the file is the header row that describes the content of every column. The field separator between columns can be any among *whitespace*, *tabulation*, *comma*, or *semicolon*. The field separator must be the same for every row in the file, including the header.

There is a minimum set of columns, that every GWAS data file should contain. In *GWAToolbox*, the following information is required for every file:

- Marker name
- Chromosome number or name
- Marker position
- Coded and non-coded allele
- Allele frequency for coded allele
- Strand
- Imputation label
- Imputation quality
- Effect estimate
- Standard error of the effect estimate
- P-value

More non-mandatory items can be included in the data file as, for example, the study sample size, the SNP call rate for genotyped SNPs, the p-value of the Hardy-Weinberg equilibrium test for genotyped SNPs, etc.

5 The Input Script

gwasqc() can analyze several GWAS data files consecutively. Instructions are provided to *gwasqc()* through a script in a text file. The format of the script file resembles the METAL input file format*.

In the input script, the user can list all GWAS file names to be analyzed and specify the format of each single GWAS file, including column names, field separator, etc. In the case when more GWAS files are in the same format, file specifications can be entered only once for all files. Example 1 illustrates the content of a hypothetical input script file.

Example 1

*<http://www.sph.umich.edu/csg/abecasis/metal/>


```

# Description of input data columns
MARKER      SNPID
CHR          Chromosome
POSITION     Position
N            n_total
ALLELE       coded_allele noncoded_allele
STRAND       strand
EFFECT       beta
STDERR       se
PVALUE       pval
FREQLABEL    allele_freq_coded_allele
IMPUTED      imputed
IMP_QUALITY  oevar_imp

# High quality filters
HQ_SNP       0.01  0.3

# Plotting filters
MAF           0.01  0.05
IMP           0.3   0.6

# Prefix for output files
PREFIX       res_

# Input file with GWA data
PROCESS       input_file.txt

```

5.1 Listing the Input Data Files

The names of the GWAS data files are specified in the input script with the command **PROCESS**[†]. If multiple files have to be checked, multiple **PROCESS** lines must be specified.

Example 2 The input script contains the following two lines:

```

PROCESS    input_file_1.txt
PROCESS    /dir_1/dir_2/input_file_2.csv

```

QC is applied first to *input_file_1.txt* and then to *input_file_2.csv*. When files reside on different directories, the full path must be specified. ◁

5.2 Describing the Input Data Columns

5.2.1 Field Separator

The field separator may be different for each GWAS data file. *gwasqc()* automatically detects the field separator of each input file *based on the first 10 rows*.

[†]*GWAToolbox* supports single line feed ('`\n`') character or carriage return character ('`\r`') followed by line feed character as the line terminators in the input files.

The user has the possibility to specify the separator manually for each GWAS file using the command **SEPARATOR**. Table 1 lists all supported separators.

Argument	Separator
COMMA	<i>comma</i>
TAB	<i>tabulation</i>
WHITESPACE	<i>whitespace</i>
SEMICOLON	<i>semicolon</i>

Table 1: The list of arguments for the SEPARATOR command.

Example 3 In the following input script:

```
PROCESS      input_file_1.txt
SEPARATOR    TAB
PROCESS      input_file_2.csv
PROCESS      input_file_3.txt
```

the field separator for the input file *input_file_1.txt* is determined automatically by *gwasqc()*; the separator for the input files *input_file_2.csv* and *input_file_3.txt* is set to tabulation by the user. ◀

If the user manually specifies the wrong field separator, then the file(s) still will be processed with this separator. As a consequence, the mandatory columns will not be detected and the user will see it in the final report.

5.2.2 Missing Values

By default, *gwasqc()* assumes that missing values are labeled as *NA*. However, the label for missing value can be specified manually by the user with the command **MISSING**.

Example 4 In the following input script:

```
MISSING      -
PROCESS      input_file_1.txt
MISSING      NA
PROCESS      input_file_2.csv
```

the *hyphen* symbol is set as symbol for missing value for the file *input_file_1.txt*. Afterwards, the coding for missing vaules is changed to *NA* and is used to process *input_file_2.csv*. ◀

5.2.3 Column Names

In table 2 the complete list of default column names for the GWAS data files is reported. These names identify uniquely the items in the GWAS data file.

Default column name(s)	Description
MARKER	Marker name
CHR	Chromosome number or name
POSITION	Marker position
ALLELE1, ALLELE2	Coded and non-coded alleles
FREQLABEL	Allele frequency for the coded allele
STRAND	Strand
IMPUTED	Label value indicating if the marker was imputed (1) or genotyped (0)
IMP_QUALITY	Imputation quality statistics; this can be different depending on the software used for imputation: MACH's <i>Rsq</i> , IMPUTE's <i>properinfo</i> , ...
EFFECT	Effect estimate
STDERR	Standard error of the effect estimate
PVALUE	P-value
HWE_PVAL	Hardy-Weinberg equilibrium p-value
CALLRATE	Genotype callrate
N	Sample size
USED_FOR_IMP	Label value indicating if a marker was used for imputation (1) or not (0)
AVPOSTPROB	Average posterior probability for imputed marker allele dosage

Table 2: The default column names.

Given that different names can be provided with the GWAS data files, *gwasqc()* allows to redefine the default values for every input file in the input script. The redefinition command consists of the default column name followed by a new column name. To redefine the default column names for *coded* and *non-coded* alleles, the command **ALLELE** is followed by two new column names.

Example 5 Let's assume to have two input files, *input_file_1.txt* and *input_file_2.txt*. In *input_file_1.txt*, the column names for the effect estimate and its standard error are *beta* and *SE*, respectively. In *input_file_2.txt*, the column name for the effect estimate is the same as in *input_file_1.txt*, but the column name for the standard error is *STDERR*. The correct column redefinitions are as follows:

```

EFFECT      beta
STDERR      SE
PROCESS     input_file_1.txt
STDERR      STDERR
PROCESS     input_file_2.csv

```

First, we redefine column names for the file *input_file_1.txt*. Notice that the column *beta* doesn't need to be redefined for file *input_file_2.csv*. However, for this file we need to redefine the column *STDERR*, returning it to its default name. ◁

Example 6 Consider an input file *input_file_1.txt* with the following names for ALLELE1 and ALLELE2: *myRefAllele* and *myNonRefAllele*. The new column definition is applied as follows:

```
ALLELE    myRefAllele myNonRefAllele
PROCESS   input_file_1.txt
```

◁

5.2.4 Case Sensitivity

By default, *gwasqc()* assumes that column names of GWAS input files are case insensitive. For example, the names *STDERR*, *StdErr*, and *STDErr* are all perfectly equivalent. This behaviour can be changed for every input file in the input script using the command **CASESENSITIVE**. Table 3 lists all possible arguments.

Argument	Description
0	Column names are case insensitive (default)
1	Column names are case sensitive

Table 3: The list of arguments for the CASESENSITIVE command.

Example 7 Consider the following commands:

```
CASESENSITIVE 1
PROCESS        input_file_1.txt
CASESENSITIVE 0
PROCESS        input_file_2.csv
```

Column names of *input_file_1.txt* are case sensitive and must correspond exactly to the default column names; column names of *input_file_2.csv* are case insensitive. ◁

5.3 Specifying Data Filters

5.3.1 Implausible Values Filters

Often, there is the necessity to identify implausible values for the statistics that will be included in the meta-analysis. Implausible values for the effect estimate, for its standard error, or the p-value are sometimes generated by the software used for the association testing. For example, in the case of a disease outcome with a small number of cases or of a SNP with very small MAF, statistical packages can report inconsistent results due to statistical algorithms that fail to converge because of data sparseness. Other types of inconsistencies can originate from errors in the file management.

In these situations, it is important to identify the SNPs with inconsistent values, so that they can be removed before starting the meta-analysis. *gwasqc()* can identify these values by using appropriate threshold values. The number of SNPs affected by this kind of problems is reported. In addition, these SNPs are

excluded from the calculation of the data quality summary statistics. Implausible value filters are used by *gwasqc()* to identify implausible values. Table 4 lists the columns for which the filters can be applied and their default thresholds.

Default column name	Default thresholds
STDERR	[0, 100000]
IMP_QUALITY	(0, 1.5)
PVALUE	(0, 1)
FREQLABEL	(0, 1)
HWE_PVAL	(0, 1)
CALLRATE	(0, 1)

Table 4: The default implausible value filters.

The default thresholds can be redefined for every column in the input script. The new thresholds for a column can be specified after the redefinition of the column name (see Section 5.2.3).

Example 8 Let's assume that the file *input_file_1.txt* has a standard error column called *STDERR* and that the corresponding column in the input file *input_file_2.csv* is called *SE*. In addition, the imputation quality column is defined as *oevar_imp* in both files. The user can redefine the column names while applying different plausibility filters:

```
STDERR      STDERR 0 80000
IMP_QUALITY  oevar_imp 0 1
PROCESS      input_file_1.txt
STDERR      SE 0 100000
PROCESS      input_file_2.csv
```

The file *input_file_1.txt* has new [0, 80000] thresholds for the standard error and new (0, 1) thresholds for the imputation quality. For the file *input_file_2.csv* the thresholds of [0, 100000] will be applied to the standard error column, while for the imputation quality column the same filters as for the *input_file_1.txt* will be applied. <

5.3.2 High Quality Filters

In many cases, analysts want to restrict the analyses to SNPs with high imputation quality and with not too small MAF. We call these SNPs '*high quality SNPs*', that is SNPs for which results should be quite robust. In the special case, when estimating the genomic inflation factor, *lambda*, to check for the presence of cryptic relatedness or hidden population sub-structures, it may be important to remove SNPs that could artificially inflate the number of significant hits. Summary statistics are calculated after excluding SNPs with low quality (*CSV* report files). Table 5 lists the default thresholds for allele frequency and imputation quality.

The default thresholds can be redefined with the command **HQ.SNP** for every input file in the input script. The command is followed by two values:

Default column name	Default thresholds
FREQLABEL	> 0.01
IMP_QUALITY	> 0.3

Table 5: The default high quality imputation filters.

the first one is the threshold for the MAF and the one is the threshold for the imputation quality.

Example 9 If we want to define 'high quality SNPs' those with MAF > 0.03 and imputation quality > 0.4, we would add the following lines to the input script:

```
HQ_SNP      0.03 0.4
PROCESS     input_file_1.txt
```

<

5.3.3 Plotting Filters

The plotting filters are used to select meaningful data to be displayed in the various summary plots. Each filter allows two threshold levels: each of them is applied dependently on the plot type and column. Figure 2 (see Section 5.4.3) shows what data and filters are used when producing plots. Table 6 lists the default threshold values.

Default column name	Default 1st level thresholds	Default 2nd level thresholds
FREQLABEL	> 0.01	> 0.05
IMP_QUALITY	> 0.3	> 0.6

Table 6: The default plotting filter.

The default threshold values for MAF and imputation quality can be redefined accordingly with the commands **MAF** and **IMP** for the every input file in the input script.

Example 10 Assume the input script contains the following commands:

```
MAF          0.02 0.03
IMP          0.3 0.5
PROCESS     input_file_1.txt
```

Here, new SNP quality thresholds are set for plotting results from *input_file_1.txt*. For the first level thresholds, we have selected MAF > 0.02 and the imputation quality > 0.3, for the second level threshold, MAF > 0.03 and imputation quality > 0.5. <

5.4 Controlling The Output

5.4.1 Output File Names

Output file names are defined based on input file names, with the addition of a specified prefix (all types of output files will share the same prefix). The prefix can be specified once for all input files, or for every single input file or groups of input files explicitly using the command **PREFIX**.

Example 11 Consider the following input script:

```
PREFIX      res_
PROCESS    input_file_1.txt
PROCESS    input_file_2.csv

PREFIX      result_
PROCESS    input_file_3.tab
```

All result output files corresponding to the input files *input_file_1.txt* and *input_file_2.csv* will be prefixed with *res_*; result output files corresponding to the input file *input_file_3.tab* will be prefixed with *result_*.

◀

5.4.2 Verbosity Level For Graphical Output

The user can control the number of output pictures with the command **VERBOSITY** (see Table 7 for the available options).

Argument	Description
1	Lowest verbosity level (default).
2	Highest verbosity level.

Table 7: The list of arguments for the VERBOSITY command.

Example 12 The input script contains the following commands:

```
VERBOSITY    2
PROCESS      input_file_1.txt
VERBOSITY    1
PROCESS      input_file_2.csv
```

The file *input_file_1.txt* is processed with the highest verbosity level and therefore all figures are generated; *input_file_2.csv* is processed with the lowest verbosity level and less output figures are generated. ◀

5.4.3 Number And Content Of Plots

Number and content of the output plots depend on the setting of the plotting filters (see Section 5.3.3) and on the available information in the input files. Figure 2 shows all the dependencies. If some dependency is not satisfied because of missing data or filter setting, then the corresponding plot is not produced or may be truncated at different levels.

Plots	Columns												
	PVALUE	EFFECT	STDERR	FREQLABEL	HWE_PVAL	CALLRATE	N	IMP_QUALITY	Q2_MAF	Q2_IMP	Q2_MAF_IMP	BOXPLOTS	BOXPLOTS_HQ
PVALUE	●								●	●	●		
EFFECT		●							●	●		●	●
STDERR	●	●	●						●	●	●	●	●
FREQLABEL	●	●		●					●		●		●
HWE_PVAL					●								
CALLRATE						●							
N							●						
IMP_QUALITY	●	●						●		●	●	●	●
Filters													
MAF level 1	●	●							●		●		●
MAF level 2									●		●		
IMP level 1	●	●								●	●	●	●
IMP level 2										●	●		

Figure 2: The dependencies of graphical outputs on columns and filters.

When multiple files are processed, boxplots from the distributions of the effect estimates are displayed in a single graph for across-study comparison. It is possible to specify the width of each box based on one of the other available information (typically the sample size). As an argument, **BOXPLOTWIDTH** requires one of the default column names. If **BOXPLOTWIDTH** is not specified all boxplots have the same width.

It is also possible to define labels for each input file, to be used in the plots instead of the full file names, which could be too long and, therefore, clutter the plots.

Example 13 Let *n_total* be the column name which identifies the sample size in the input file *input_file_1.txt*, and *samplesize* the corresponding name in *input_file_2.csv*. Consider the following input script:

```

N          n_total
PROCESS    input_file_1.txt first
N          samplesize
PROCESS    /dir_1/dir_2/input_file_2.csv second
BOXPLOTWIDTH  N

```

The width of the first boxplot for the input file *input_file_1.txt* depends on the *n_total* column, while the width of the second boxplot for the input file *input_file_2.csv* depends on the *samplesize* column. The labels "first" and "second" will be used to label the two studies in the plots. ◁

6 The Output Files

The typical output of the *gwasqc()* function consists of the following files:

1. **Graphical files** (*PNG* file extension) include **histograms** and **boxplots** of the distribution of the main statistics from each GWAS file: effect estimates, imputation quality index, sample size, p-value, allele frequency;

QQ-plots of the p-value distribution are also provided to investigate the presence of study-design bias. See Figure 2 for an exhaustive list of available plots.

2. **Textual report** (*TXT* file extension) contains information on the GWAS file-format quality and statistics summarizing the distribution of all data present in the GWAS files (effect estimates, p-values, etc.). The statistics provided in this report can be compared with the graphical output described above.
3. **Comma separated report** (*CSV* file extension) contains summary statistics for the high quality SNPs, as they have been defined by the user when setting the parameter **HQ-SNP** in the QC script (see Section 5.3.2). The tabular format of this file is intended to be useful for users who wants to perform additional analyses and compare results from different GWAS files without having to manage the large original GWAS files. Figure 3 shows an example of such file. For each of the variables listed by

	N	EFFECT	STDERR	PVALUE	MAF
N	2543887	2542617	2542422	2542617	2543887
N_HQ	2448544	2448544	2448544	2448544	2448544
N_NAs	0	1270	1270	1270	0
Mean	516	-0.00137	0.295076	0.500823	0.230416
StdDev	0	0.333379	0.158119	0.28873	0.143312
Min	516	-8.76038	0.17353	7.34E-08	0.010001
Max	516	6.12165	2.53341	1	0.5
Median	516	-0.00149	0.233707	0.502629	0.217008
Skewness	nan	-0.06933	2.886758	-0.00935	0.214034
Kurtosis	nan	9.109873	11.14145	-1.20085	-1.18454

	IMPUTED	IMP_QUALITY	CALLRATE	HWE_PVAL	STD_EFFECT_0.5
N	2543887	2543887	2543887	NA	NA
N_HQ	2448544	2448544	2448544	NA	NA
N_NAs	0	0	0	NA	NA
Mean	1	0.958868	1	NA	NA
StdDev	0	0.097734	0	NA	NA
Min	1	0.300006	1	NA	NA
Max	1	1.0467	1	NA	NA
Median	1	0.994577	1	NA	NA
Skewness	nan	-3.90226	nan	NA	-0.00587
Kurtosis	nan	16.76683	nan	NA	3.297394

	STD_EFFECT_0.75	STD_EFFECT_0.95	STD_EFFECT_0.99	STD_EFFECT_1
N	NA	NA	NA	NA
N_HQ	NA	NA	NA	NA
N_NAs	NA	NA	NA	NA
Mean	NA	NA	NA	NA
StdDev	NA	NA	NA	NA
Min	NA	NA	NA	NA
Max	NA	NA	NA	NA
Median	NA	NA	NA	NA
Skewness	-0.01036	0.022903	0.040024	-0.07186
Kurtosis	3.822136	4.712606	5.861049	9.066722

Figure 3: The comma separated report file.

columns, summary statistics are reported by row. Specifically:

- **N** – number of SNPs with available information (i.e. non-missing values).
- **N_HQ** – number of high quality SNPs with information available (i.e. non-missing values).
- **N_NAs** – number of SNPs with missing values for the specific field of interest.
- **Mean**, **StdDev**, **Min**, **Max**, **Median**, **Skewness**, and **Kurtosis** are referred to the distribution of the specific field.

Columns include:

- **N** – study sample size.
 - **EFFECT**, **STDERR**, and **PVALUE** – summaries of the SNP-phenotype associations.
 - **MAF**, **IMPUTED**, **IMP_QUALITY**, **CALLRATE**, and **HWE_PVAL** – summaries of the genotype distribution and quality.
 - **STD_EFFECT_0.5**, **STD_EFFECT_0.75**, **STD_EFFECT_0.95**, **STD_EFFECT_0.99**, **STD_EFFECT_1** – these columns are only of interest for the reported skewness and kurtosis of the standardized effect estimates (beta/SE). The numbers 0.5, 0.75, 0.95, 0.99, and 1 are referred to the percentage of SNPs with the highest p-values chosen to study the effect size distribution (see Section 8.1 for additional details).
4. **An HTML report** (*main.html* file) combines both textual and graphical output, allowing the user to easily surfing across the results and across the studies included in the QC analysis. In operating systems allowing graphical interfaces, the *HTML* document should be the first file to be opened to investigate the results.

7 Example

This is an embedded R code example. All input files for this example are located in the subdirectory *doc* of the installed *GWAtoolbox* package.

Consider the following five GWAS data files: *gwa_data_example_1.txt*, *gwa_data_example_2.tbl*, *gwa_data_example_3.csv*, *gwa_data_example_4.txt* and *gwa_data_example_5.csv*. The first file contains 16 whitespace-separated columns:

```
> t <- read.table("gwa_data_example_1.txt", header = T, nrow = 1,
+               sep = " ")
> colnames(t)

[1] "SNPID"           "chr"             "position"        "coded_all"
[5] "noncoded_all"    "strand_genome"   "beta"            "SE"
[9] "pval"            "AF_coded_all"    "n_total"         "oevar_imp"
[13] "avpostprob"      "callrate"        "HWE_pval"        "used_for_imp"
[17] "imputed"
```

The second file contains 16 tab-separated columns:

```
> t <- read.table("gwa_data_example_2.tbl", header = T, nrow = 1,
+               sep = "\t")
> colnames(t)

[1] "SNPID"           "chr"             "position"        "coded_all"
[5] "noncoded_all"    "strand_genome"   "beta"            "StdErr"
[9] "p"               "AF_coded_all"    "n_total"         "oevar_imp"
[13] "avpostprob"      "callrate"        "HWE_pval"        "used_for_imp"
[17] "imputed"
```

Analogously, we can preview the headers of the other three files. To perform the QC of these files we prepare a simple input script: *GWAS_script.txt*. Below are listed the commands which were included in the script:

```
> cat(readLines("GWASQC_script.txt"), sep = "\n")

# Column names
ALLELE                coded_all noncoded_all
CALLRATE              callrate
CHR                   chr
EFFECT                 beta
FREQLABEL             AF_coded_all
HWE_PVAL              HWE_pval
IMPUTED               imputed
IMP_QUALITY           oever_imp
MARKER                SNPID
N                     n_total
POSITION              position
PVALUE                pval
STRAND                strand_genome
STDERR               SE
USED_FOR_IMP          used_for_imp

# Plotting filters for the MAF and imputation quality
MAF 0.01    0.05
IMP 0.3     0.5

# Prefix for output files
PREFIX      gwasqc_

# Column N controls the width of boxplots
BOXPLOTWIDTH N

# Input file and its short name for labels
PROCESS gwa_data_example_1.txt Study1

PVALUE      p
STDERR      StdErr

PROCESS gwa_data_example_2.tbl Study2
PROCESS gwa_data_example_3.csv Study3

PVALUE      pvalue

PROCESS gwa_data_example_4.txt Study4
PROCESS gwa_data_example_5.csv Study5
```

When the input script is ready, load the *GWAtoolbox* library and call the *gwasqc()* function as follows:

```
> library(GWAtoolbox)
> gwasqc("GWASQC_script.txt")
```

8 Between-study comparisons

8.1 Comparing skewness and kurtosis of effect size distribution

Association of genetic markers with continuous or binary phenotypes is generally assessed by the use of linear models, where an effect estimate and its standard error are used to summarize the evidence of the association. The effect estimate is usually represented by the beta coefficient of the linear regression model. For binary outcome, this correspond to the log(odds ratio) obtained from logistic regression models. Let's define with θ our parameter of interest, and with $SE(\theta)$ its standard error. For large sample sizes, under the null hypothesis of no association, the distribution of $\theta/SE(\theta) \sim N(0, 1)$.

In checking the quality of GWA results, we are interested in assessing potential errors arisen during the analytical process or during the file management process. These errors could origin $\theta/SE(\theta)$ distributions that are systematically biased towards positive (or negative) values, or that are over/under-dispersed. The *kurtosis* and the *skewness* indices are the natural candidates to perform this kind of assessment. Convenient graphical display based on these two indices enables the contemporary plot and comparison of all the studies involved in the meta-analysis, with consequent identification of studies that are systematically different from each other.

Under the forms proposed by Cramer [3], the kurtosis and the skewness indices can be defined as $ku = \frac{\frac{1}{n} \sum_{i=1}^n (x_i - \bar{x})^4}{\left(\frac{1}{n} \sum_{i=1}^n (x_i - \bar{x})^2\right)^2} - 3$ and $sk = \frac{\frac{1}{n} \sum_{i=1}^n (x_i - \bar{x})^3}{\left(\frac{1}{n} \sum_{i=1}^n (x_i - \bar{x})^2\right)^{\frac{3}{2}}}$.

The skewness index assesses the symmetry of a distribution around its central value, and the kurtosis index assesses the dispersion of the distribution around its central value. If $\theta/SE(\theta) \sim N(0, 1)$, then for large sample sizes, $sk_\theta = sk(\theta/SE(\theta)) \rightarrow 0$ and $ku_\theta = ku(\theta/SE(\theta)) \rightarrow 0$ (Fisher [4]; Joanes and Gill [5]).

In a GWAS setting, we can assume that the 50% of SNPs with largest p-value are not associated with the phenotype of interest and so, they can be used to represent the null situation. Notice that the 50% of SNPs with worst p-values correspond to the concept of the genomic control inflation factor, which is estimated based on the median chi-square distribution from the p-values.

In real world applications, distribution of $\theta/SE(\theta)$ for 50% worst SNPs is not normally distributed, because SNPs are not independent each other. This situation is more pronounced in presence of genotype imputation, which causes an excess of effect estimates that are close to 0. For this reason, it is not realistic to expect that $ku_\theta \rightarrow 0$: the distribution is leptokuric, and so $ku_\theta > 0$. However, as long as the studies involved in the meta-analysis used similar imputation reference platforms, ku_θ should be similar for all studies. For what concerns the skewness, there is no good reason why sk_θ of the 50% worst SNPs shouldn't approximate 0.

Then, for given K studies, we can estimate $ku_{\theta,k}$ and $sk_{\theta,k}$ for $k = 1, \dots, K$, and we can plot the two vectors $sk_{\theta,[1,\dots,K]}$ vs. $ku_{\theta,[1,\dots,K]}$ in a Cartesian diagram, with every point representing a different study. We expect all studies to cluster around the same point at $sk = 0$, with similar kurtosis values. Studies that show strong departures from the main cluster could be submitted to detailed investigation in order to detect the reason of such discrepancy. In general,

departures along the sk -axis are more serious than departures on the sk -axis, because the first ones will introduce systematic bias in the meta-analysis with one or a few studies that are systematically different from the others, with effect estimates that are more often in one direction.

This diagnostic plot is shown in Figure 4 where the set of SNPs with 50% largest p-values is compared with other sets of SNPs with the largest 75%, 95%, and 100% (i.e. all SNPs are considered) p-values, respectively. Highlighted are points (studies) that are largely different from the ones in the cluster. In the last scenario, all SNPs are included and the bias is given by the SNPs that are truly associated with the phenotype. Outlier studies, that have been identified in the 50% scatterplot, are colored in red also in the other situations for comparison.




Figure 4: The skewness and kurtosis plot.

The *GWAtoolbox* allows automatic comparison of skewness and kurtosis of effect size distribution between GWA studies. The *gwasqc()* function estimates the skewness and kurtosis statistics during the QC workflow and includes them into the *CSV* reports. Then, the auxiliary *kusk_check()* function can be used to export this information to a *R* data frame and to produce diagnostic plots. As input, it requires the same script used for *gwasqc()* and assumes that all the *CSV* reports are located in the current working directory. An optional list consisting of any of integer number among 50, 75, 95, 99, and 100, can be specified: numbers correspond to the percentage of SNPs to be considered as representing the null distribution.

Example 14 We report the commands to obtain the scatterplot shown in Figure 4, when the 50% SNPs with largest p-values is considered.

```

> W <- kusk_check("GWASQC_script.txt", worst = c(50), plot = TRUE)
> points(W$sk50[W$ku50 > 5], W$ku50[W$ku50 > 5], pch = 22, bg = 2,
+       cex = 2)
> text(W$sk50[W$ku50 > 5], W$ku50[W$ku50 > 5], labels = W$study[W$ku50 >
+       5], cex = 1, pos = 4)
<

```

Currently, no automatic method to identify outlier studies is implemented and that the user needs to define his/her own criteria for the outlier identification.

8.2 Precision of the effect estimates by sample size

A different graphical test, that allows the comparison of studies against each other, is based on the assessment of the distribution of estimates' precision vs. study sample size. In general, the average $SE(\theta)$ is expected to be inversely proportional to the study sample size. The auxiliary *dispersion_check()* function plots a scatterplot of the $mean(SE(\theta))$ vs. the median sample size of all studies, as depicted in the example Figure 5. Over-dispersion is defined as the presence of larger SEs than expected given the study sample size and under-dispersion is meant to be the opposite phenomenon. For example, a study with unmodeled relatedness or population stratification may present SEs that are smaller than another study of similar sample size where these issues were accounted for properly.




Figure 5: Schematic representation of the dispersion plot and its interpretation.

The *dispersion_check()* function uses the *CSV* reports generated by *gwasqc()*. As input it requires the same script as *gwasqc()* and assumes that all the *CSV* reports files are located in the current working directory. If the study sample size is missing from GWAS files, then the function has an optional parameter allowing to specify a vector with all study sample sizes. The function returns an *R* data frame with the information extracted from the *CSV* reports and produces the diagnostic plot.

Example 15

```
> Z <- dispersion_check("GWASQC_script.txt", plot = TRUE)
> Z
```

	study	mean_se	median_n
1	Study1	0.01188491	1201
2	Study2	0.03206312	201
3	Study3	0.01270057	1000
4	Study4	0.02030638	437
5	Study5	0.01532013	721

```
> text(Z$median_n, Z$mean_se, labels = Z$study, pos = c(2, 4, 2,
+ 1, 2))

> Z <- dispersion_check("GWASQC_script.txt", sample_sizes = c(1200,
+ 200, 1000, 500, 700), plot = TRUE)
> Z
```

	study	mean_se	median_n
1	Study1	0.01188491	1200
2	Study2	0.03206312	200
3	Study3	0.01270057	1000
4	Study4	0.02030638	500
5	Study5	0.01532013	700

```
> text(Z$median_n, Z$mean_se, labels = Z$study, pos = c(2, 4, 2,
+ 1, 2))
<
```

9 Additional Tools

In addition to the facilities for the data QC, *GWAtoolbox* also includes facilities including routines for the data manipulation after the GWAS meta-analysis has been done and functions that facilitates result annotation. These extensions are currently being debugged and will be made available soon.

References

- [1] Cristen J. Willer, Yun Li, and Gonçalo R. Abecasis. (2010) **METAL: fast and efficient meta-analysis of genomewide association scans**. *Bioinformatics* 26: 2190-2191.

- [2] Paul I.W. de Bakker, Manuel A.R. Ferreira, Xiaoming Jia, Benjamin M. Neale, Soumya Raychaudhuri, and Benjamin F. Voight (2008) **Practical aspects of imputation-driven meta-analysis of genome-wide association studies**. Hum. Mol. Genet. 17: R122-R128.
- [3] H. Cramer (1946) **Mathematical Methods of Statistics**. Princeton: Princeton University Press.
- [4] R.A. Fisher (1930) **The moments of the distribution for normal samples of measures of departure from normality**. Proc. R. Soc. Series A 130:16-28.
- [5] D.N. Joanes, C.A. Gill (1998) **Comparing Measures of Sample Skewness and Kurtosis**. J. Royal Stat. Soc. Series D (The Statistician) 47(1):183-189.

Index

ALLELE, [11](#)
ALLELE1, [11](#)
ALLELE2, [11](#)
AVPOSTPROB, [11](#)

BOXPLOTWIDTH, [16](#)

CALLRATE, [11](#), [13](#)
CASESENSITIVE, [12](#)
CHR, [11](#)
COMMA, [10](#)

EFFECT, [11](#)

FREQLABEL, [11](#), [13](#), [14](#)

HQ_SNP, [13](#), [17](#)
HWE_PVAL, [11](#), [13](#)

IMP, [14](#)
IMP_QUALITY, [11](#), [13](#), [14](#)
IMPUTED, [11](#)

MAF, [14](#)
MARKER, [11](#)
MISSING, [10](#)

N, [11](#)

POSITION, [11](#)
PREFIX, [15](#)
PROCESS, [9](#)
PVALUE, [11](#), [13](#)

SEMICOLON, [10](#)
SEPARATOR, [10](#)
STDERR, [11](#), [13](#)
STRAND, [11](#)

TAB, [10](#)

USED_FOR_IMP, [11](#)

VERBOSITY, [15](#)

WHITESPACE, [10](#)